Machine Learning Mastery
Making Developers Awesome at Machine Learning

Search...   🔍

A Gentle Introduction to Cross-Entropy for Machine Learning

by **Jason Brownlee** on October 21, 2019 in **Probability**

Last Updated on December 22, 2020

Cross-entropy is commonly used in machine learning as a loss function.

Cross-entropy is a measure from the field of information theory, building upon entropy and generally calculating the difference between two probability distributions. It is closely related to but is different from KL divergence that calculates the relative entropy between two probability distributions, whereas cross-entropy can be thought to calculate the total entropy between the distributions.

Cross-entropy is also related to and often confused with logistic loss, called log loss. Although the two measures are derived from a different source, when used as loss functions for classification models, both measures calculate the same quantity and can be used interchangeably.

In this tutorial, you will discover cross-entropy for machine learning.

After completing this tutorial, you will know:

- How to calculate cross-entropy from scratch and using standard machine learning libraries.
- Cross-entropy can be used as a loss function when optimizing classification models like logistic regression and artificial neural networks.
- Cross-entropy is different from KL divergence but can be calculated using KL divergence, and is different from log loss but calculates the same quantity when used as a loss function.

Kick-start your project with my new book Probability for Machine Learning, including *step-by-step tutorials* and the *Python source code* files for all examples.

Let's get started.

- **Update Oct/2019**: Gave an example of cross-entropy for identical distributions and updated description for this case (thanks Ron U). Added an example of calculating the entropy of the known class labels.
- **Update Nov/2019**: Improved structure and added more explanation of entropy. Added intuition for predicted class probabilities.
- **Update Dec/2020**: Tweaked the introduction to information and entropy to be clearer.

A Gentle Introduction to Cross-Entropy for Machine Learning
Photo by Jerome Bon, some rights reserved.

Tutorial Overview

This tutorial is divided into five parts; they are:

1. What Is Cross-Entropy?
2. Cross-Entropy Versus KL Divergence
3. How to Calculate Cross-Entropy
    1. Two Discrete Probability Distributions
    2. Calculate Cross-Entropy Between Distributions
    3. Calculate Cross-Entropy Between a Distribution and Itself
    4. Calculate Cross-Entropy Using KL Divergence
4. Cross-Entropy as a Loss Function
    1. Calculate Entropy for Class Labels
    2. Calculate Cross-Entropy Between Class Labels and Probabilities
    3. Calculate Cross-Entropy Using Keras
    4. Intuition for Cross-Entropy on Predicted Probabilities
5. Cross-Entropy Versus Log Loss
    1. Log Loss is the Negative Log Likelihood
    2. Log Loss and Cross Entropy Calculate the Same Thing

What Is Cross-Entropy?

Cross-entropy is a measure of the difference between two probability distributions for a given random variable or set of events.

You might recall that **information** quantifies the number of bits required to encode and transmit an event. Lower probability events have more information, higher probability events have less information.

In information theory, we like to describe the "*surprise*" of an event. An event is more surprising the less likely it is, meaning it contains more information.

- **Low Probability Event** (*surprising*): More information.
- **Higher Probability Event** (*unsurprising*): Less information.

Information $h(x)$ can be calculated for an event x, given the probability of the event $P(x)$ as follows:

- $h(x) = -\log(P(x))$

Entropy is the number of bits required to transmit a randomly selected event from a probability distribution. A skewed distribution has a low entropy, whereas a distribution where events have equal probability has a larger entropy.

A skewed probability distribution has less "surprise" and in turn a low entropy because likely events dominate. Balanced distribution are more surprising and turn have higher entropy because events are equally likely.

- **Skewed Probability Distribution** (*unsurprising*): Low entropy.
- **Balanced Probability Distribution** (*surprising*): High entropy.

Entropy *H(x)* can be calculated for a random variable with a set of *x* in *X* discrete states discrete states and their probability *P(x)* as follows:

- H(X) = – sum x in X P(x) * log(P(x))

If you would like to know more about calculating information for events and entropy for distributions see this tutorial:

- A Gentle Introduction to Information Entropy

Cross-entropy builds upon the idea of entropy from information theory and calculates the number of bits required to represent or transmit an average event from one distribution compared to another distribution.

> *… the cross entropy is the average number of bits needed to encode data coming from a source with distribution p when we use model q …*

— Page 57, Machine Learning: A Probabilistic Perspective, 2012.

The intuition for this definition comes if we consider a target or underlying probability distribution P and an approximation of the target distribution Q, then the cross-entropy of Q from P is the number of additional bits to represent an event using Q instead of P.

The cross-entropy between two probability distributions, such as Q from P, can be stated formally as:

- H(P, Q)

Where H() is the cross-entropy function, P may be the target distribution and Q is the approximation of the target distribution.

Cross-entropy can be calculated using the probabilities of the events from P and Q, as follows:

- H(P, Q) = – sum x in X P(x) * log(Q(x))

Where P(x) is the probability of the event x in P, Q(x) is the probability of event x in Q and log is the base-2 logarithm, meaning that the results are in bits. If the base-e or natural logarithm is used

instead, the result will have the units called nats.

This calculation is for discrete probability distributions, although a similar calculation can be used for continuous probability distributions using the integral across the events instead of the sum.

The result will be a positive number measured in bits and will be equal to the entropy of the distribution if the two probability distributions are identical.

Note: this notation looks a lot like the joint probability, or more specifically, the joint entropy between P and Q. This is misleading as we are scoring the difference between probability distributions with cross-entropy. Whereas, joint entropy is a different concept that uses the same notation and instead calculates the uncertainty across two (or more) random variables.

Want to Learn Probability for Machine Learning

Take my free 7-day email crash course now (with sample code).

Click to sign-up and also get a free PDF Ebook version of the course.

Download Your FREE Mini-Course

Cross-Entropy Versus KL Divergence

Cross-entropy is not KL Divergence.

Cross-entropy is related to divergence measures, such as the Kullback-Leibler, or KL, Divergence that quantifies how much one distribution differs from another.

Specifically, the KL divergence measures a very similar quantity to cross-entropy. It measures the average number of extra bits required to represent a message with Q instead of P, not the total number of bits.

> *In other words, the KL divergence is the average number of extra bits needed to encode the data, due to the fact that we used distribution q to encode the data instead of the true distribution p.*

— Page 58, Machine Learning: A Probabilistic Perspective, 2012.

As such, the KL divergence is often referred to as the "*relative entropy*."

- **Cross-Entropy**: Average number of total bits to represent an event from Q instead of P.
- **Relative Entropy** (*KL Divergence*): Average number of extra bits to represent an event from Q instead of P.

KL divergence can be calculated as the negative sum of probability of each event in P multiples by the log of the probability of the event in Q over the probability of the event in P. Typically, log base-2 so that the result is measured in bits.

- KL(P || Q) = – sum x in X P(x) * log(Q(x) / P(x))

The value within the sum is the divergence for a given event.

As such, we can calculate the cross-entropy by adding the entropy of the distribution plus the additional entropy calculated by the KL divergence. This is intuitive, given the definition of both calculations; for example:

- H(P, Q) = H(P) + KL(P || Q)

Where H(P, Q) is the cross-entropy of Q from P, H(P) is the entropy of P and KL(P || Q) is the divergence of Q from P.

Entropy can be calculated for a probability distribution as the negative sum of the probability for each event multiplied by the log of the probability for the event, where log is base-2 to ensure the result is in bits.

- H(P) = – sum x on X p(x) * log(p(x))

Like KL divergence, cross-entropy is not symmetrical, meaning that:

- H(P, Q) != H(Q, P)

As we will see later, both cross-entropy and KL divergence calculate the same quantity when they are used as loss functions for optimizing a classification predictive model. It is under this context that you might sometimes see that cross-entropy and KL divergence are the same.

For a lot more detail on the KL Divergence, see the tutorial:

- How to Calculate the KL Divergence for Machine Learning

How to Calculate Cross-Entropy

In this section we will make the calculation of cross-entropy concrete with a small example.

Two Discrete Probability Distributions

Consider a random variable with three discrete events as different colors: red, green, and blue.

We may have two different probability distributions for this variable; for example:

```
1  ...
2  # define distributions
3  events = ['red', 'green', 'blue']
4  p = [0.10, 0.40, 0.50]
5  q = [0.80, 0.15, 0.05]
```

We can plot a bar chart of these probabilities to compare them directly as probability histograms.

The complete example is listed below.

```
1   # plot of distributions
2   from matplotlib import pyplot
3   # define distributions
4   events = ['red', 'green', 'blue']
5   p = [0.10, 0.40, 0.50]
6   q = [0.80, 0.15, 0.05]
7   print('P=%.3f Q=%.3f' % (sum(p), sum(q)))
8   # plot first distribution
9   pyplot.subplot(2,1,1)
10  pyplot.bar(events, p)
11  # plot second distribution
12  pyplot.subplot(2,1,2)
13  pyplot.bar(events, q)
14  # show the plot
15  pyplot.show()
```

Running the example creates a histogram for each probability distribution, allowing the probabilities for each event to be directly compared.

We can see that indeed the distributions are different.

Histogram of Two Different Probability Distributions for the Same Random Variable

Calculate Cross-Entropy Between Distributions

Next, we can develop a function to calculate the cross-entropy between the two distributions.

We will use log base-2 to ensure the result has units in bits.

```
1  # calculate cross entropy
2  def cross_entropy(p, q):
3      return -sum([p[i]*log2(q[i]) for i in range(len(p))])
```

We can then use this function to calculate the cross-entropy of P from Q, as well as the reverse, Q from P.

```
1  ...
2  # calculate cross entropy H(P, Q)
3  ce_pq = cross_entropy(p, q)
4  print('H(P, Q): %.3f bits' % ce_pq)
5  # calculate cross entropy H(Q, P)
6  ce_qp = cross_entropy(q, p)
7  print('H(Q, P): %.3f bits' % ce_qp)
```

Tying this all together, the complete example is listed below.

```
1  # example of calculating cross entropy
2  from math import log2
3
4  # calculate cross entropy
5  def cross_entropy(p, q):
6      return -sum([p[i]*log2(q[i]) for i in range(len(p))])
7
8  # define data
9  p = [0.10, 0.40, 0.50]
10 q = [0.80, 0.15, 0.05]
11 # calculate cross entropy H(P, Q)
12 ce_pq = cross_entropy(p, q)
13 print('H(P, Q): %.3f bits' % ce_pq)
14 # calculate cross entropy H(Q, P)
15 ce_qp = cross_entropy(q, p)
16 print('H(Q, P): %.3f bits' % ce_qp)
```

Running the example first calculates the cross-entropy of Q from P as just over 3 bits, then P from Q as just under 3 bits.

```
1  H(P, Q): 3.288 bits
2  H(Q, P): 2.906 bits
```

Calculate Cross-Entropy Between a Distribution and Itself

If two probability distributions are the same, then the cross-entropy between them will be the entropy of the distribution.

We can demonstrate this by calculating the cross-entropy of P vs P and Q vs Q.

The complete example is listed below.

```
1  # example of calculating cross entropy for identical distributions
2  from math import log2
3
4  # calculate cross entropy
5  def cross_entropy(p, q):
6      return -sum([p[i]*log2(q[i]) for i in range(len(p))])
7
8  # define data
9  p = [0.10, 0.40, 0.50]
10 q = [0.80, 0.15, 0.05]
11 # calculate cross entropy H(P, P)
12 ce_pp = cross_entropy(p, p)
13 print('H(P, P): %.3f bits' % ce_pp)
14 # calculate cross entropy H(Q, Q)
15 ce_qq = cross_entropy(q, q)
16 print('H(Q, Q): %.3f bits' % ce_qq)
```

Running the example first calculates the cross-entropy of Q vs Q which is calculated as the entropy for Q, and P vs P which is calculated as the entropy for P.

```
1  H(P, P): 1.361 bits
2  H(Q, Q): 0.884 bits
```

Calculate Cross-Entropy Using KL Divergence

We can also calculate the cross-entropy using the KL divergence.

The cross-entropy calculated with KL divergence should be identical, and it may be interesting to calculate the KL divergence between the distributions as well to see the relative entropy or additional bits required instead of the total bits calculated by the cross-entropy.

First, we can define a function to calculate the KL divergence between the distributions using log base-2 to ensure the result is also in bits.

```
1  # calculate the kl divergence KL(P || Q)
2  def kl_divergence(p, q):
3      return sum(p[i] * log2(p[i]/q[i]) for i in range(len(p)))
```

Next, we can define a function to calculate the entropy for a given probability distribution.

```
1  # calculate entropy H(P)
2  def entropy(p):
3      return -sum([p[i] * log2(p[i]) for i in range(len(p))])
```

Finally, we can calculate the cross-entropy using the *entropy()* and *kl_divergence()* functions.

```
1  # calculate cross entropy H(P, Q)
2  def cross_entropy(p, q):
3      return entropy(p) + kl_divergence(p, q)
```

To keep the example simple, we can compare the cross-entropy for H(P, Q) to the KL divergence KL(P || Q) and the entropy H(P).

The complete example is listed below.

```
1   # example of calculating cross entropy with kl divergence
2   from math import log2
3
4   # calculate the kl divergence KL(P || Q)
5   def kl_divergence(p, q):
6       return sum(p[i] * log2(p[i]/q[i]) for i in range(len(p)))
7
8   # calculate entropy H(P)
9   def entropy(p):
10      return -sum([p[i] * log2(p[i]) for i in range(len(p))])
11
12  # calculate cross entropy H(P, Q)
13  def cross_entropy(p, q):
14      return entropy(p) + kl_divergence(p, q)
15
16  # define data
17  p = [0.10, 0.40, 0.50]
18  q = [0.80, 0.15, 0.05]
```

```
19  # calculate H(P)
20  en_p = entropy(p)
21  print('H(P): %.3f bits' % en_p)
22  # calculate kl divergence KL(P || Q)
23  kl_pq = kl_divergence(p, q)
24  print('KL(P || Q): %.3f bits' % kl_pq)
25  # calculate cross entropy H(P, Q)
26  ce_pq = cross_entropy(p, q)
27  print('H(P, Q): %.3f bits' % ce_pq)
```

Running the example, we can see that the cross-entropy score of 3.288 bits is comprised of the entropy of P 1.361 and the additional 1.927 bits calculated by the KL divergence.

This is a useful example that clearly illustrates the relationship between all three calculations.

```
1  H(P): 1.361 bits
2  KL(P || Q): 1.927 bits
3  H(P, Q): 3.288 bits
```

Cross-Entropy as a Loss Function

Cross-entropy is widely used as a loss function when optimizing classification models.

Two examples that you may encounter include the logistic regression algorithm (a linear classification algorithm), and artificial neural networks that can be used for classification tasks.

> " … using the cross-entropy error function instead of the sum-of-squares for a classification problem leads to faster training as well as improved generalization.

— Page 235, Pattern Recognition and Machine Learning, 2006.

**Never miss a tutorial:**

Classification problems are those that involve one or more input variables and the prediction of a class label.

Classification tasks that have just two labels for the output variable are referred to as binary classification problems, whereas those problems with more than two labels are referred to as categorical or multi-class classification problems.

Picked for you:

How to Use ROC Curves and Precision-Recall Curves for Classification in Python

- **Binary Classification**: Task of predicting one of two class labels for a given example.
- **Multi-Class Classification**: Task of predicting one of more than two class labels for a given example.

How and When to Use a Calibrated Classification Model with scikit-learn

We can see that the idea of cross-entropy may be useful for optimizing a classification model.

How to Implement Bayesian Optimization from Scratch in Python

Each example has a known class label with a probability of 1.0, and a probability of 0.0 for all other labels. A model can estimate the probability of an example belonging to each class label. Cross-

can then be used to calculate the difference between the two probability distributions.

As such, we can map the classification of one example onto the idea of a random variable with a probability distribution as follows:

- **Random Variable**: The example for which we require a predicted class label.
- **Events**: Each class label that could be predicted.

In classification tasks, we know the target probability distribution P for an input as the class label 0 or 1 interpreted as probabilities as "*impossible*" or "*certain*" respectively. These probabilities have no surprise at all, therefore they have no information content or zero entropy.

Our model seeks to approximate the target probability distribution Q.

In the language of classification, these are the actual and the predicted probabilities, or *y* and *yhat*.

- **Expected Probability** (*y*): The known probability of each class label for an example in the dataset (P).
- **Predicted Probability** (*yhat*): The probability of each class label an example predicted by the model (Q).

We can, therefore, estimate the cross-entropy for a single prediction using the cross-entropy calculation described above; for example.

- H(P, Q) = – sum x in X P(x) * log(Q(x))

Where each *x* in *X* is a class label that could be assigned to the example, and *P(x)* will be 1 for the known label and 0 for all other labels.

The cross-entropy for a single example in a binary classification task can be stated by unrolling the sum operation as follows:

- H(P, Q) = – (P(class0) * log(Q(class0)) + P(class1) * log(Q(class1)))

You may see this form of calculating cross-entropy cited in textbooks.

If there are just two class labels, the probability is modeled as the Bernoulli distribution for the positive class label. This means that the probability for class 1 is predicted by the model directly, and the probability for class 0 is given as one minus the predicted probability, for example:

- **Predicted P(class0)** = 1 – yhat
- **Predicted P(class1)** = yhat

When calculating cross-entropy for classification tasks, the base-e or natural logarithm is used. This

means that the units are in nats, not bits.

We are often interested in minimizing the cross-entropy for the model across the entire training dataset. This is calculated by calculating the average cross-entropy across all training examples.

Calculate Entropy for Class Labels

Recall that when two distributions are identical, the cross-entropy between them is equal to the entropy for the probability distribution.

Class labels are encoded using the values 0 and 1 when preparing data for classification tasks.

For example, if a classification problem has three classes, and an example has a label for the first class, then the probability distribution will be [1, 0, 0]. If an example has a label for the second class, it will have a probability distribution for the two events as [0, 1, 0]. This is called a one hot encoding.

This probability distribution has no information as the outcome is certain. We know the class. Therefore the entropy for this variable is zero.

This is an important concept and we can demonstrate it with a worked example.

Pretend with have a classification problem with 3 classes, and we have one example that belongs to each class. We can represent each example as a discrete probability distribution with a 1.0 probability for the class to which the example belongs and a 0.0 probability for all other classes.

We can calculate the entropy of the probability distribution for each "*variable*" across the "*events*".

The complete example is listed below.

```
1  # entropy of examples from a classification task with 3 classes
2  from math import log2
3  from numpy import asarray
4
5  # calculate entropy
6  def entropy(p):
7      return -sum([p[i] * log2(p[i]) for i in range(len(p))])
8
9  # class 1
10 p = asarray([1,0,0]) + 1e-15
11 print(entropy(p))
12 # class 2
13 p = asarray([0,1,0]) + 1e-15
14 print(entropy(p))
15 # class 3
16 p = asarray([0,0,1]) + 1e-15
17 print(entropy(p))
```

Running the example calculates the entropy for each random variable.

We can see that in each case, the entropy is 0.0 (actually a number very close to zero).

Note that we had to add a very small value to the 0.0 values to avoid the *log()* from blowing up, as we cannot calculate the log of 0.0.

```
1  9.805612959471341e-14
2  9.805612959471341e-14
3  9.805612959471341e-14
```

As such, the entropy of a known class label is always 0.0.

This means that the cross entropy of two distributions (real and predicted) that have the same probability distribution for a class label, will also always be 0.0.

Recall that when evaluating a model using cross-entropy on a training dataset that we average the cross-entropy across all examples in the dataset.

Therefore, a cross-entropy of 0.0 when training a model indicates that the predicted class probabilities are identical to the probabilities in the training dataset, e.g. zero loss.

We could just as easily minimize the KL divergence as a loss function instead of the cross-entropy.

Recall that the KL divergence is the extra bits required to transmit one variable compared to another. It is the cross-entropy without the entropy of the class label, which we know would be zero anyway.

As such, minimizing the KL divergence and the cross entropy for a classification task are identical.

> *Minimizing this KL divergence corresponds exactly to minimizing the cross-entropy between the distributions.*

— Page 132, Deep Learning, 2016.

In practice, a cross-entropy loss of 0.0 often indicates that the model has overfit the training dataset, but that is another story.

Calculate Cross-Entropy Between Class Labels and Probabilities

The use of cross-entropy for classification often gives different specific names based on the number of classes, mirroring the name of the classification task; for example:

- **Binary Cross-Entropy**: Cross-entropy as a loss function for a binary classification task.
- **Categorical Cross-Entropy**: Cross-entropy as a loss function for a multi-class classification task.

We can make the use of cross-entropy as a loss function concrete with a worked example.

Consider a two-class classification task with the following 10 actual class labels (P) and predicted class labels (Q).

```
1  ...
2  # define classification data
3  p = [1, 1, 1, 1, 1, 0, 0, 0, 0, 0]
4  q = [0.8, 0.9, 0.9, 0.6, 0.8, 0.1, 0.4, 0.2, 0.1, 0.3]
```

We can enumerate these probabilities and calculate the cross-entropy for each using the cross-entropy function developed in the previous section using *log()* (natural logarithm) instead of *log2()*.

```
1  # calculate cross entropy
2  def cross_entropy(p, q):
3      return -sum([p[i]*log(q[i]) for i in range(len(p))])
```

For each actual and predicted probability, we must convert the prediction into a distribution of probabilities across each event, in this case, the classes {0, 1} as 1 minus the probability for class 0 and probability for class 1.

We can then calculate the cross-entropy and repeat the process for all examples.

```
1   ...
2   # calculate cross entropy for each example
3   results = list()
4   for i in range(len(p)):
5       # create the distribution for each event {0, 1}
6       expected = [1.0 - p[i], p[i]]
7       predicted = [1.0 - q[i], q[i]]
8       # calculate cross entropy for the two events
9       ce = cross_entropy(expected, predicted)
10      print('>[y=%.1f, yhat=%.1f] ce: %.3f nats' % (p[i], q[i], ce))
11      results.append(ce)
```

Finally, we can calculate the average cross-entropy across the dataset and report it as the cross-entropy loss for the model on the dataset.

```
1  ...
2  # calculate the average cross entropy
3  mean_ce = mean(results)
4  print('Average Cross Entropy: %.3f nats' % mean_ce)
```

Tying this all together, the complete example is listed below.

```
1  # calculate cross entropy for classification problem
2  from math import log
3  from numpy import mean
4
5  # calculate cross entropy
6  def cross_entropy(p, q):
7      return -sum([p[i]*log(q[i]) for i in range(len(p))])
8
9  # define classification data
```

```
10  p = [1, 1, 1, 1, 1, 0, 0, 0, 0, 0]
11  q = [0.8, 0.9, 0.9, 0.6, 0.8, 0.1, 0.4, 0.2, 0.1, 0.3]
12  # calculate cross entropy for each example
13  results = list()
14  for i in range(len(p)):
15      # create the distribution for each event {0, 1}
16      expected = [1.0 - p[i], p[i]]
17      predicted = [1.0 - q[i], q[i]]
18      # calculate cross entropy for the two events
19      ce = cross_entropy(expected, predicted)
20      print('>[y=%.1f, yhat=%.1f] ce: %.3f nats' % (p[i], q[i], ce))
21      results.append(ce)
22
23  # calculate the average cross entropy
24  mean_ce = mean(results)
25  print('Average Cross Entropy: %.3f nats' % mean_ce)
```

Running the example prints the actual and predicted probabilities for each example and the cross-entropy in nats.

The final average cross-entropy loss across all examples is reported, in this case, as 0.247 nats.

```
1   >[y=1.0, yhat=0.8] ce: 0.223 nats
2   >[y=1.0, yhat=0.9] ce: 0.105 nats
3   >[y=1.0, yhat=0.9] ce: 0.105 nats
4   >[y=1.0, yhat=0.6] ce: 0.511 nats
5   >[y=1.0, yhat=0.8] ce: 0.223 nats
6   >[y=0.0, yhat=0.1] ce: 0.105 nats
7   >[y=0.0, yhat=0.4] ce: 0.511 nats
8   >[y=0.0, yhat=0.2] ce: 0.223 nats
9   >[y=0.0, yhat=0.1] ce: 0.105 nats
10  >[y=0.0, yhat=0.3] ce: 0.357 nats
11  Average Cross Entropy: 0.247 nats
```

This is how cross-entropy loss is calculated when optimizing a logistic regression model or a neural network model under a cross-entropy loss function.

Calculate Cross-Entropy Using Keras

We can confirm the same calculation by using the *binary_crossentropy()* function from the Keras deep learning API to calculate the cross-entropy loss for our small dataset.

The complete example is listed below.

Note: This example assumes that you have the Keras library installed (e.g. version 2.3 or higher) and configured with a backend library such as TensorFlow (version 2.0 or higher). If not, you can skip running this example.

```
1  # calculate cross entropy with keras
2  from numpy import asarray
3  from keras import backend
4  from keras.losses import binary_crossentropy
5  # prepare classification data
6  p = asarray([1, 1, 1, 1, 1, 0, 0, 0, 0, 0])
```

```
7  q = asarray([0.8, 0.9, 0.9, 0.6, 0.8, 0.1, 0.4, 0.2, 0.1, 0.3])
8  # convert to keras variables
9  y_true = backend.variable(p)
10 y_pred = backend.variable(q)
11 # calculate the average cross-entropy
12 mean_ce = backend.eval(binary_crossentropy(y_true, y_pred))
13 print('Average Cross Entropy: %.3f nats' % mean_ce)
```

Running the example, we can see that the same average cross-entropy loss of 0.247 nats is reported.

This confirms the correct manual calculation of cross-entropy.

```
1  Average Cross Entropy: 0.247 nats
```

Intuition for Cross-Entropy on Predicted Probabilities

We can further develop the intuition for the cross-entropy for predicted class probabilities.

For example, given that an average cross-entropy loss of 0.0 is a perfect model, what do average cross-entropy values greater than zero mean exactly?

We can explore this question no a binary classification problem where the class labels as 0 and 1. This is a discrete probability distribution with two events and a certain probability for one event and an impossible probability for the other event.

We can then calculate the cross entropy for different "*predicted*" probability distributions transitioning from a perfect match of the target distribution to the exact opposite probability distribution.

We would expect that as the predicted probability distribution diverges further from the target distribution that the cross-entropy calculated will increase.

The example below implements this and plots the cross-entropy result for the predicted probability distribution compared to the target of [0, 1] for two events as we would see for the cross-entropy in a binary classification task.

```
1  # cross-entropy for predicted probability distribution vs label
2  from math import log
3  from matplotlib import pyplot
4
5  # calculate cross-entropy
6  def cross_entropy(p, q, ets=1e-15):
7      return -sum([p[i]*log(q[i]+ets) for i in range(len(p))])
8
9  # define the target distribution for two events
10 target = [0.0, 1.0]
11 # define probabilities for the first event
12 probs = [1.0, 0.9, 0.8, 0.7, 0.6, 0.5, 0.4, 0.3, 0.2, 0.1, 0.0]
13 # create probability distributions for the two events
```

```
14 dists = [[1.0 - p, p] for p in probs]
15 # calculate cross-entropy for each distribution
16 ents = [cross_entropy(target, d) for d in dists]
17 # plot probability distribution vs cross-entropy
18 pyplot.plot([1-p for p in probs], ents, marker='.')
19 pyplot.title('Probability Distribution vs Cross-Entropy')
20 pyplot.xticks([1-p for p in probs], ['[%.1f,%.1f]'%(d[0],d[1]) for d in dists], rotation=
21 pyplot.subplots_adjust(bottom=0.2)
22 pyplot.xlabel('Probability Distribution')
23 pyplot.ylabel('Cross-Entropy (nats)')
24 pyplot.show()
```

Running the example calculates the cross-entropy score for each probability distribution then plots the results as a line plot.

We can see that as expected, cross-entropy starts at 0.0 (far left point) when the predicted probability distribution matches the target distribution, then steadily increases as the predicted probability distribution diverges.

We can also see a dramatic leap in cross-entropy when the predicted probability distribution is the exact opposite of the target distribution, that is, [1, 0] compared to the target of [0, 1].

Line Plot of Probability Distribution vs Cross-Entropy for a Binary Classification Task

We are not going to have a model that predicts the exact opposite probability distribution for all cases on a binary classification task.

As such, we can remove this case and re-calculate the plot.

The updated version of the code is listed below.

```python
# cross-entropy for predicted probability distribution vs label
from math import log
from matplotlib import pyplot

# calculate cross-entropy
def cross_entropy(p, q, ets=1e-15):
	return -sum([p[i]*log(q[i]+ets) for i in range(len(p))])

# define the target distribution for two events
target = [0.0, 1.0]
# define probabilities for the first event
probs = [1.0, 0.9, 0.8, 0.7, 0.6, 0.5, 0.4, 0.3, 0.2, 0.1]
# create probability distributions for the two events
dists = [[1.0 - p, p] for p in probs]
# calculate cross-entropy for each distribution
ents = [cross_entropy(target, d) for d in dists]
# plot probability distribution vs cross-entropy
pyplot.plot([1-p for p in probs], ents, marker='.')
pyplot.title('Probability Distribution vs Cross-Entropy')
pyplot.xticks([1-p for p in probs], ['[%.1f,%.1f]'%(d[0],d[1]) for d in dists], rotation=
pyplot.subplots_adjust(bottom=0.2)
pyplot.xlabel('Probability Distribution')
pyplot.ylabel('Cross-Entropy (nats)')
pyplot.show()
```

Running the example gives a much better idea of the relationship between the divergence in probability distribution and the calculated cross-entropy.

We can see a super-linear relationship where the more the predicted probability distribution diverges from the target, the larger the increase in cross-entropy.

Line Plot of Probability Distribution vs Cross-Entropy for a Binary Classification Task With Extreme Case Removed

A plot like this can be used as a guide for interpreting the average cross-entropy reported for a model for a binary classification dataset.

For example, you can use these cross-entropy values to interpret the mean cross-entropy reported by Keras for a neural network model on a binary classification task, or a binary classification model in scikit-learn evaluated using the logloss metric.

You can use it to answer the general question:

> **What is a good cross-entropy score?**

If you are working in nats (and you usually are) and you are getting mean cross-entropy less than 0.2, you are off to a good start, and less than 0.1 or 0.05 is even better.

On the other hand, if you are getting mean cross-entropy greater than 0.2 or 0.3 you can probably improve, and if you are getting a mean cross-entropy greater than 1.0, then something is going on

and you're making poor probability predictions on many examples in your dataset.

We can summarise these intuitions for the mean cross-entropy as follows:

- **Cross-Entropy = 0.00**: Perfect probabilities.
- **Cross-Entropy < 0.02**: Great probabilities.
- **Cross-Entropy < 0.05**: On the right track.
- **Cross-Entropy < 0.20**: Fine.
- **Cross-Entropy > 0.30**: Not great.
- **Cross-Entropy > 1.00**: Terrible.
- **Cross-Entropy > 2.00** Something is broken.

This listing will provide a useful guide when interpreting a cross-entropy (log loss) from your logistic regression model, or your artificial neural network model.

You can also calculate separate mean cross-entropy scores per-class and help tease out on which classes you're model has good probabilities, and which it might be messing up.

Cross-Entropy Versus Log Loss

Cross-Entropy is not Log Loss, but they calculate the same quantity when used as loss functions for classification problems.

Log Loss is the Negative Log Likelihood

Logistic loss refers to the loss function commonly used to optimize a logistic regression model.

It may also be referred to as logarithmic loss (which is confusing) or simply log loss.

Many models are optimized under a probabilistic framework called the maximum likelihood estimation, or MLE, that involves finding a set of parameters that best explain the observed data.

This involves selecting a likelihood function that defines how likely a set of observations (data) are given model parameters. When a log likelihood function is used (which is common), it is often referred to as optimizing the log likelihood for the model. Because it is more common to minimize a function than to maximize it in practice, the log likelihood function is inverted by adding a negative sign to the front. This transforms it into a Negative Log Likelihood function or NLL for short.

In deriving the log likelihood function under a framework of maximum likelihood estimation for Bernoulli probability distribution functions (two classes), the calculation comes out to be:

- negative log-likelihood(P, Q) = -(P(class0) * log(Q(class0)) + P(class1) * log(Q(class1)))

This quantity can be averaged over all training examples by calculating the average of the log of the likelihood function.

Negative log-likelihood for binary classification problems is often shortened to simply "log loss" as the loss function derived for logistic regression.

- log loss = negative log-likelihood, under a Bernoulli probability distribution

We can see that the negative log-likelihood is the same calculation as is used for the cross-entropy for Bernoulli probability distribution functions (two events or classes). In fact, the negative log-likelihood for Multinoulli distributions (multi-class classification) also matches the calculation for cross-entropy.

For more on log loss and the negative log likelihood, see the tutorial:

- A Gentle Introduction to Logistic Regression With Maximum Likelihood Estimation

Log Loss and Cross Entropy Calculate the Same Thing

For classification problems, "*log loss*", "*cross-entropy*" and "*negative log-likelihood*" are used interchangeably.

More generally, the terms "*cross-entropy*" and "*negative log-likelihood*" are used interchangeably in the context of loss functions for classification models.

> *The negative log-likelihood for logistic regression is given by [...] This is also called the cross-entropy error function.*

— Page 246, Machine Learning: A Probabilistic Perspective, 2012.

Therefore, calculating log loss will give the same quantity as calculating the cross-entropy for Bernoulli probability distribution. We can confirm this by calculating the log loss using the log_loss() function from the scikit-learn API.

Calculating the average log loss on the same set of actual and predicted probabilities from the previous section should give the same result as calculating the average cross-entropy.

The complete example is listed below.

```
1  # calculate log loss for classification problem with scikit-learn
2  from sklearn.metrics import log_loss
3  from numpy import asarray
4  # define classification data
5  p = [1, 1, 1, 1, 1, 0, 0, 0, 0, 0]
6  q = [0.8, 0.9, 0.9, 0.6, 0.8, 0.1, 0.4, 0.2, 0.1, 0.3]
```

```
 7  # define data as expected, e.g. probability for each event {0, 1}
 8  y_true = asarray([[1-v, v] for v in p])
 9  y_pred = asarray([[1-v, v] for v in q])
10  # calculate the average log loss
11  ll = log_loss(y_true, y_pred)
12  print('Average Log Loss: %.3f' % ll)
```

Running the example gives the expected result of 0.247 log loss, which matches 0.247 nats when calculated using the average cross-entropy.

```
1  Average Log Loss: 0.247
```

This does not mean that log loss calculates cross-entropy or cross-entropy calculates log loss.

Instead, they are different quantities, arrived at from different fields of study, that under the conditions of calculating a loss function for a classification task, result in an equivalent calculation and result. Specifically, a cross-entropy loss function is equivalent to a maximum likelihood function under a Bernoulli or Multinoulli probability distribution.

This demonstrates a connection between the study of maximum likelihood estimation and information theory for discrete probability distributions.

It is not limited to discrete probability distributions, and this fact is surprising to many practitioners that hear it for the first time.

Specifically, a linear regression optimized under the maximum likelihood estimation framework assumes a Gaussian continuous probability distribution for the target variable and involves minimizing the mean squared error function. This is equivalent to the cross-entropy for a random variable with a Gaussian probability distribution.

> *Any loss consisting of a negative log-likelihood is a cross-entropy between the empirical distribution defined by the training set and the probability distribution defined by model. For example, mean squared error is the cross-entropy between the empirical distribution and a Gaussian model.*

— Page 132, Deep Learning, 2016.

This is a little mind blowing, and comes from the field of differential entropy for continuous random variables.

It means that if you calculate the mean squared error between two Gaussian random variables that cover the same events (have the same mean and standard deviation), then you are calculating the cross-entropy between the variables.

It also means that if you are using mean squared error loss to optimize your neural network model

for a regression problem, you are in effect using a cross entropy loss.

Further Reading

This section provides more resources on the topic if you are looking to go deeper.

Tutorials

- A Gentle Introduction to Information Entropy
- How to Calculate the KL Divergence for Machine Learning
- A Gentle Introduction to Logistic Regression With Maximum Likelihood Estimation
- How to Choose Loss Functions When Training Deep Learning Neural Networks
- Loss and Loss Functions for Training Deep Learning Neural Networks

Books

- Machine Learning: A Probabilistic Perspective, 2012.
- Pattern Recognition and Machine Learning, 2006.
- Deep Learning, 2016.

API

- Usage of loss functions, Keras API.
- sklearn.metrics.log_loss API.

Articles

- Kullback-Leibler divergence, Wikipedia.
- Cross entropy, Wikipedia.
- Joint Entropy, Wikipedia.
- Loss functions for classification, Wikipedia.
- Differential entropy, Wikipedia.

Summary

In this tutorial, you discovered cross-entropy for machine learning.

Specifically, you learned:

- How to calculate cross-entropy from scratch and using standard machine learning libraries.
- Cross-entropy can be used as a loss function when optimizing classification models like logistic regression and artificial neural networks.
- Cross-entropy is different from KL divergence but can be calculated using KL divergence, and is different from log loss but calculates the same quantity when used as a loss function.

Do you have any questions?
Ask your questions in the comments below and I will do my best to answer.

Get a Handle on Probability for Machine Learning!

Develop Your Understanding of Probability

...with just a few lines of python code

Discover how in my new Ebook:
[Probability for Machine Learning](#)

It provides **self-study tutorials** and **end-to-end projects** on:
Bayes Theorem, *Bayesian Optimization*, *Distributions*, *Maximum Likelihood*,
Cross-Entropy, *Calibrating Models*
and much more...

Finally Harness Uncertainty in Your Projects

Skip the Academics. Just Results.

<div>SEE WHAT'S INSIDE</div>

Tweet Tweet Share Share

More On This Topic

Information Gain and Mutual Information for Machine Learning

How to Choose Loss Functions When Training Deep…


A Gentle Introduction to Information Entropy

Loss and Loss Functions for Training Deep Learning…

Nested Cross-Validation for Machine Learning with Python

Repeated k-Fold Cross-Validation for Model…

About Jason Brownlee

Jason Brownlee, PhD is a machine learning specialist who teaches developers how to get results with modern machine learning methods via hands-on tutorials.

View all posts by Jason Brownlee →

64 Responses to *A Gentle Introduction to Cross-Entropy for Machine Learning*

Markus October 22, 2019 at 6:57 am #

REPLY ↩

Hi

Thanks for this blog post.

What confuses me a bit is the fact that we interpret the labels 0 and 1 in the example as the probability values for calculating the cross entropy between the target distribution and the predicted distribution!

What if the labels were 4 and 7 instead of 0 and 1?!

Jason Brownlee October 22, 2019 at 1:44 pm #

REPLY ↩

Thanks.

Good question. For binary classification we map the labels, whatever they are to 0 and 1.

I recommend reading about the Bernoulli distribution:

https://machinelearningmastery.com/discrete-probability-distributions-for-machine-learning/

sharpblade4 October 22, 2019 at 5:35 pm #

REPLY ↩

Hi Jason,
A small fix suggestion: in the beginning of the article in section "What Is Cross-Entropy?" you've mentioned that "The result will be a positive number measured in bits and 0 if the two probability

distributions are identical.".

However, the cross entropy for the same probability-distributions H(P,P) is the entropy for the probability-distribution H(P), opposed to KL divergence of the same probability-distribution which would indeed outcome zero.

Thanks,
Ron U

Jason Brownlee October 23, 2019 at 6:34 am #

Thanks Ron!

I'll schedule time to update the post and give an example of exactly what you're referring to. E.g.:

```
1   # example of calculating cross entropy
2   from math import log2
3
4   # calculate cross entropy
5   def cross_entropy(p, q):
6       return -sum([p[i]*log2(q[i]) for i in range(len(p))])
7
8   # define data
9   p = [0.10, 0.40, 0.50]
10  q = [0.80, 0.15, 0.05]
11  # calculate cross entropy H(P, P)
12  ce_pp = cross_entropy(p, p)
13  print('H(P, P): %.3f bits' % ce_pp)
14  # calculate cross entropy H(Q, Q)
15  ce_qq = cross_entropy(q, q)
16  print('H(Q, Q): %.3f bits' % ce_qq)
```

H(P, P): 1.361 bits
H(Q, Q): 0.884 bits

Update: I have updated the post to correctly discuss this case.

Thanks Again!

Hugh Brown October 26, 2019 at 4:39 pm #

Why not use `zip`?

```
1   # calculate cross entropy
2   def cross_entropy(p, q):
3       return -sum(pp*log2(qq) for pp, qq in zip(p, q))
```

Jason Brownlee October 27, 2019 at 5:43 am #

Thanks for the tip Hugh, that is a much cleaner approach!

Anthony The Koala October 28, 2019 at 3:48 am #

Dear Dr Jason,

In the last few lines under the subheading "How to Calculate Cross-Entropy", you had the simple example with the following outputs:

```
1  H(P): 1.361 bits          #entropy
2  KL(P || Q): 1.927 bits #kl divergence
3  H(P, Q): 3.288 bits       #cross entropy = entropy + kl divergence
```

What is the interpretation of these figures in 'plain English' please.

For example if the above example produced the following result:

```
1  H(P): 0.361 bits          #entropy
2  KL(P || Q): 2.927 bits #kl divergence
3  H(P, Q): 3.288 bits       #cross entropy = entropy + kl divergence
```

Or if we had this:

```
1  H(P): 1.927 bits          #entropy
2  KL(P || Q): 0.361 bits #kl divergence
3  H(P, Q): 3.288 bits       #cross entropy = entropy + kl divergence
```

Here is another example of made up figures.

```
1  H(P) = 0.05 bits         #entropy
2  KL(P||Q) = 0.2 bits   #kl diverengence
3  H(P,Q) = 0.25 bits     #cross entropy = entropy + kl divergence.
```

Comparing the first output to the 'made up figures' does the lower the number of bits mean a better fit? A more predictable model?

Also:

I understand that a bit is a base 2 number. Eg 1 = 1(base 10), 11 = 3 (base 10), 101 = 5 (base 10). The number of bits in a base 2 system is an integer. What does a fraction of bit mean?

Anthony of Sydney

Jason Brownlee October 28, 2019 at 6:10 am #

Typically we use cross-entropy to evaluate a model, e.g. true classes vs probability predictions.

In that case would compare the average cross-entropy calculated across all examples and a lower value would represent a better fit. Interpreting the specific figures is often not useful.

Anthony The Koala October 28, 2019 at 6:18 am #

Dear Dr Jason,
Thank you for response.
The other question please:
How can you have a fraction of a bit. For example entropy = 3.2285 bits. What is 0.2285 bits.
Thank you,
Anthony of Sydney

Jason Brownlee October 28, 2019 at 7:00 am #

Recall, it is an average over a distribution with many events.

Think of it more of a measure and less like the crisp bits in a computer.

Farhan October 28, 2019 at 6:32 pm #

Great Article, Hope to see more more content on machine learning and AI.

Jason Brownlee October 29, 2019 at 5:21 am #

Thanks!

Confused October 31, 2019 at 1:26 pm #

Hi,

could you provide an example of this sentence "The entropy for a distribution of all 0 or all 1 values or mixtures of these values will equal 0.0."?

Does this mean a distribution with a mixture of these values, eg. the distribution with P(X=1) = 0.4 and P(X=0) = 0.6 has entropy zero? But this should not be the case because 0.4 * log(0.4) + 0.6 * log(0.6) is not zero.

Thank you!

Jason Brownlee October 31, 2019 at 1:46 pm #

Good question.

Sorry that is confusing.

I mean that the probability distribution for a class label will always be zero.

I have updated the tutorial to be clearer and given a worked example.

Xin November 13, 2019 at 9:35 am #

Hi Jason! This is the best article I've ever seen on cross entropy and KL-divergence! Finally I can understand them 😉 Thank you so much for the comprehensive article.

I have one small question: in the secion "Intuition for Cross-Entropy on Predicted Probabilities", in the first code block to plot the visualization, the code is as follows:

```
# define the target distribution for two events
target = [0.0, 0.1]
# define probabilities for the first event
probs = [0.0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0]
# create probability distributions for the two events
dists = [[p, 1.0 – p] for p in probs]
# calculate cross-entropy for each distribution
ents = [cross_entropy(target, d) for d in dists]
```

I do not quite understand why the target probability for the two events are [0.0, 0.1]? Could you explain a bit more? Thank you!

Jason Brownlee November 13, 2019 at 1:45 pm #

Thanks! I worked really hard on it and I'm so happy that it's appreciated 🙂

Yes, looks like a typo. I'll fix it ASAP. It should be [0,1].

Update: I have updated the code and re-generated the plots. Thanks again!

Scott December 14, 2019 at 8:36 am #

Jason, I so appreciate all your various posts on ML topics. If I may add one comment regarding what I've found helpful in the past:

One point that I didn't see really emphasized here that I've seen in other treatments (e.g., https://tdhopper.com/blog/cross-entropy-and-kl-divergence/) is that cross-entropy and KL difference "differ by a constant", i.e. in your expression

H(P, Q) = H(P) + KL(P || Q),

the H(P) is constant with respect to Q. In most ML tasks, P is usually fixed as the "true" distribution" and Q is the distribution we are iteratively trying to refine until it matches P.

"In many of these situations, p is treated as the 'true' distribution, and q as the model that we're trying to optimize…. Because p is fixed, $H(p)$ doesn't change with the parameters of the model, and can be disregarded in the loss function." (https://stats.stackexchange.com/questions/265966/why-do-we-use-kullback-leibler-divergence-rather-than-cross-entropy-in-the-t-sne/265989)

You do get to this when you say "As such, minimizing the KL divergence and the cross entropy for a classification task are identical."

And yet for me at least, knowing that the two "differ by a constant" makes it intuitively obvious why minimizing one is the same as minimizing the other, even if they're actually intended to measure different things.

…Thus I think this "differ by a constant" is another reason that people get mixed up about cross-entropy vs KL divergence, and why guides like yours are so helpful!

Jason Brownlee December 15, 2019 at 6:01 am #

Thanks for your note Scott.

Yes, H(P) is the entropy of the distribution. This becomes 0 when class labels are 0 and 1. I outline this at the end of the post when we talk about class labels.

A constant of 0 in that case means using KL divergence and cross entropy result in the same numbers, e.g. the kl divergence.

Arij February 11, 2020 at 6:52 am #

thanks for a grate article!
but what confused me that in your article you have mentioned that
"

The result will be a positive number measured in bits and will be equal to the entropy of the distribution if the two probability distributions are identical."

then again mentioned

"If two probability distributions are the same, then the cross-entropy between them will be the entropy of the distribution."

then

"This means that the cross entropy of two distributions (real and predicted) that have the same probability distribution for a class label, will also always be 0.0."

"Therefore, a cross-entropy of 0.0 when training a model indicates that the predicted class probabilities are identical to the probabilities in the training dataset, e.g. zero loss."

Am I missing something?

REPLY ↩

Jason Brownlee February 11, 2020 at 1:42 pm #

The statements are correct.

The cross-entropy will be the entropy between the distributions if the distributions are identical.

In this case, if we are working with class labels like 0 and 1, then the entropy for two identical distributions will be zero.

We demonstrate this with a worked example in the above tutorial.

Does that help?

REPLY ↩

Arij February 11, 2020 at 5:36 pm #

it was not about examples, they were understandable, thanks.
Just I could not imagine and understand them numerically.
Reading them again I understand that when the values of any distribution are only one or zero then entropy, cross-entropy, KL all will be zero.

they will have values just in case they have values between 0 and 1 also

REPLY ↩

Jason Brownlee February 12, 2020 at 5:43 am #

Exactly!

REPLY ↩

Grzegorz Kępisty April 17, 2020 at 4:23 pm #

Good morning Jason!

Thank you for great post!

Question on KL Divergence: In its definition we have log2(p[i]/q[i]) which suggests a possibility of zero division error. Is it a probable issue in real applications? Or for some reason it does not occur? Regards!

REPLY ↩

Jason Brownlee April 18, 2020 at 5:41 am #

Good question, no problem as probabilities are always greater than zero, so log never blows up.

More on kl divergence here too:
https://machinelearningmastery.com/divergence-between-probability-distributions/

REPLY ↩

Allan April 27, 2020 at 2:25 am #

Is it possible to use KL divergence as a classification criterion?

REPLY ↩

Jason Brownlee April 27, 2020 at 5:37 am #

Probably, it would be the same as log loss and cross entropy when using class labels instead of probabilities.

REPLY ↩

Allan April 27, 2020 at 12:46 pm #

Thanks for your reply. So let say the final calculation result is "Average Log Loss", what does this value implies meaning?

REPLY ↩

Jason Brownlee April 28, 2020 at 6:39 am #

Average difference between the probability distributions of expected and

predicted values in bits.

Gledson April 29, 2020 at 6:21 am #

Hello Jason, Congratulations on the explanation. How are you? I have a doubt. If I have log(0), I get -Inf on my crossentropy. Should I replace -Inf with some value? If so, what value?

Best,
Gledson.

Jason Brownlee April 29, 2020 at 6:36 am #

You cannot log a zero. It is a good idea to always add a tiny value to anything to log, e.g. log(value + 1e-8)

Zeinab May 13, 2020 at 3:47 am #

Hi Jason,

Thanks for all your great post, I've read some of them.

I'm working on traffic classification and I've converted my data to string of bits, I want to use cross-entropy on bytes.

Assume below lists:

p = [1, 0, 1, 1, 0, 0, 1, 0]
q = [1, 1, 1, 0, 1, 0, 0, 1]

When I use `–sum([p[i] * log2(q[i]) for i in range(len(p))])`, I encounter this error :ValueError: math domain error. S

Would you please tell me what I'm doing wrong here and how can I implement cross-entropy on a list of bits?

Bests,

Jason Brownlee May 13, 2020 at 6:43 am #

Can't calculate log of 0.0. Try adding a tiny value to the equation, e.g. 1e-8 or 1e-15

Zeinab May 13, 2020 at 11:03 pm #

REPLY ↩

I've converted the traffic to string of bits, it's not just some random numbers that I can add any value.

TuanVu May 13, 2020 at 7:58 pm #

REPLY ↩

Why we use log function for cross entropy?

Jason Brownlee May 14, 2020 at 5:46 am #

REPLY ↩

Good question, perhaps start here:

https://machinelearningmastery.com/what-is-information-entropy/

TuanVu May 14, 2020 at 9:22 am #

REPLY ↩

thanks you !!

Jason Brownlee May 14, 2020 at 1:25 pm #

REPLY ↩

You're welcome.

Zahir May 29, 2020 at 3:36 pm #

REPLY ↩

Hello Jason,
Thank you so much for all your great posts.
I have a quesion, if we have conditional entropy $H(y|x) = -sum\ P(x,y)\ log(P(y|x))$
could we say that it is equal to cross-entropy $H(x,y) = -\ sum\ y\ log\ y\hat{}$?
And if that correct where we could say that? i.e., under what assumptions.

Thank you in advance.

Jason Brownlee May 30, 2020 at 5:52 am #

I think you're asking me if the conditional entropy is the same as the cross entropy. I don't think it is off the cuff, but perhaps confirm with a good textbook.

Zahir May 30, 2020 at 10:23 am #

Thank you so much for your replay,
I found it in "Privacy-Preserving Adversarial Networks" paper, the authors get a conditional entropy as a cost function, but when they implement the article, they use cross-entropy. But they don't say why?

Jason Brownlee May 31, 2020 at 6:17 am #

Fascinating.

Perhaps email the authors directly.

BARURI SAI AVINASH November 1, 2020 at 4:22 pm #

How can be Number of bits per charecter in text generation is equal to loss ???

Jason Brownlee November 2, 2020 at 6:38 am #

Perhaps try re-reading the above tutorial that lays it all out.

Eric Orr December 22, 2020 at 11:05 am #

This is excellent Introduction to Cross-Entropy. It seems that one of the following sentences may have a typo in the stated notion of "surprise". My first impression is that the second sentence should have said "are less surprising". Is that true?

"Low probability events are more surprising therefore have a larger amount of information. Whereas probability distributions where the events are equally likely are more surprising and have larger entropy."

Jason Brownlee December 22, 2020 at 1:34 pm #

Thanks.

Yes it could be clearer. Information is about events, entropy is about distributions, cross-entropy is about comparing distributions.

Surprise means something different when talking about information/events as compared to entropy/distributions.

I mixed the discussion of the two at the start of the tutorial. I have updated the text to be clearer.

Also see this:
https://machinelearningmastery.com/what-is-information-entropy/

Does that help?

Eric Orr December 22, 2020 at 11:48 am #

After additional consideration, it appears that the second sentence might instead be related as follows.

"In probability distributions where the events are equally likely, no events have larger or smaller likelihood (smaller or larger surprise, respectively), and the distribution has larger entropy."

Sorry for belaboring this. It is a good point but sometimes confusing.

stefani January 6, 2021 at 11:19 pm #

They say that you have understood a concept when you can describe it with very simple words and I feel that is the case here. Wonderful job!

Jason Brownlee January 7, 2021 at 6:18 am #

Thanks!

Aaron February 20, 2021 at 8:52 am #

"Relative Entropy (KL Divergence): Average number of extra bits to represent an event from Q instead of P."

Souldn't it rather say: Relative Entropy (KL Divergence): Average number of extra bits to represent an event from P using Q instead of P.

Since the expectation is over P(x)*[…]

dimid June 1, 2021 at 8:18 pm #

REPLY ↰

Dear Jason,

Thank you for this clear explanation. Why is it the case that "using the cross-entropy error function instead of the sum-of-squares for a classification problem leads to faster training as well as improved generalization"?

Jason Brownlee June 2, 2021 at 5:42 am #

REPLY ↰

It is a better match for the objective of the optimization problem.

ajay June 18, 2021 at 4:05 pm #

REPLY ↰

thank you for your good explanation. i am training a neural network is this result considerable for

dataset accuracy cross entropy loss
train 97% 0.07
val 96.5% 0.48
test 98% 0.13

Jason Brownlee June 19, 2021 at 5:46 am #

REPLY ↰

Perhaps focus on the metric you want to optimize, rather than the loss.

sarah September 28, 2021 at 3:50 pm #

REPLY ↰

Hi Jason,

can you advise on how to utilise correntropy for regression problem optimization, please ?

Adrian Tam September 29, 2021 at 11:52 pm # REPLY ↰

I don't think cross entropy applies to regression. How would you define the probability distribution required for that?

sarah September 30, 2021 at 11:01 am # REPLY ↰

the correntropy (a special case of cross entropy) is a nonlinear and local similarity metric measuring the similarity between two random variables in a neighbourhood of joint space.

C. Liangjun, P. Honeine, Q. Hua, Z. Jihong, and S. Xia, "Correntropy- based robust multilayer extreme learning machines," Pattern Recognit., vol. 84, pp. 357–370, Dec. 2018.

Adrian Tam October 1, 2021 at 12:36 pm # REPLY ↰

Thanks for pointing this out. I thought that was a typo. Looking at equations (13) and (14) of the cited paper, you should see it is a modified version of MSE. The problem is you need to create your own function (I don't see correntropy is defined in scikit-learn, for example) and not every model allows you to plug-in your own loss function. Probably the linear regression library in scikit-learn can't do it and you need to resort to scipy to do the regression manually.

sara October 12, 2021 at 3:05 pm # REPLY ↰

I see

Thanks Adiran

nkm December 13, 2021 at 4:33 am # REPLY ↰

Hello Dr. Adrian,
thanks for your valuable to give a deeper understanding. During implementation, I am using label smoothing (loss=tensorflow.keras.losses.CategoricalCrossentropy(label_smoothing=0.2)) to improve

performance as a regularizer. I have observed that though performance improves quite a good way, the training and validation loss remains very high (near 0.5), though performance learning curves are converging well. Without label smoothing, losses go very low (near 0.1) but do not converge well or show overfitting.

I would like to know:

1. Is this high loss of near 0.5 with good fit is acceptable?
2. Can I reduce this loss anyway?
3. Do you suggest an alternative to label smoothing ?

Note: I have observed in all blogs label smoothing in an increase of this loss.

Kindly guide.

Adrian Tam December 15, 2021 at 6:02 am #   REPLY ↩

Loss is a continuous measure. So for classification tasks, we should also look at accuracy. Check that out and see if it gives you a better idea.

---

nkm December 15, 2021 at 10:24 pm #   REPLY ↩

Thanks for the guidance. I will work on it.

Adrian Tam December 17, 2021 at 7:07 am #   REPLY ↩

You're welcomed.

---

milad December 24, 2021 at 8:44 pm #   REPLY ↩

hi
i'm working on cross entropy and i want to calculate two c e between two path ;
the code auther wrote on his paper is
CrossEntropy
Input:
Sequence X, Sequence Y
Output:
e: cross entropy of X and Y

```
begin
for (int n=1; n<5; n++)
begin
for each (ngram in X.ngrams(n)) //reads each ngram with length n within X
begin
(token, history) = split (ngram);//splits ngram so that ngram = history + token
a = X.P(ngram);//probability ofngram in X sequence
b = Y.ConditionalP(token,history);//probability of tokengiven history in Y sequence
logb = (b!=0 ? Log(b) : Log(1/X.ngramCount));
c = a × logb;
e += c;
end
end
end
```

can you help me write this in python??

James Carmichael February 27, 2022 at 12:43 pm #   REPLY ↩

Thanks for asking.

I'm eager to help, but I just don't have the capacity to debug code for you.

I am happy to make some suggestions:

Consider aggressively cutting the code back to the minimum required. This will help you isolate the problem and focus on it.
Consider cutting the problem back to just one or a few simple examples.
Consider finding other similar code examples that do work and slowly modify them to meet your needs. This might expose your misstep.
Consider posting your question and code to StackOverflow.

Leave a Reply

Name (required)

Email (will not be published) (required)

SUBMIT COMMENT

Welcome!

I'm *Jason Brownlee* PhD
and I **help developers** get results with **machine learning**.

Read more

LinkedIn | Twitter | Facebook | Newsletter | RSS

Privacy | Disclaimer | Terms | Contact | Sitemap | Search